AI makes us faster. Who keeps us coherent?

Governance, systems, and the new role of DesignOps in product teams



- More possibility.
- More responsibility.

The production floor is everywhere.

Product

PM + AI

- Brief
- Requirements
- Hypotheses

Research

RESEARCH + AI

- Synthesis
- Insights
- Questions

Design

DESIGN + AI

- Concepts
- Prototypes
- Patterns

Engineering

ENGINEERING + AI

- Components
- Code
- Tests

The bottleneck moves: **production → alignment → quality.**

Julian Alvarado

Xbox Design Ops

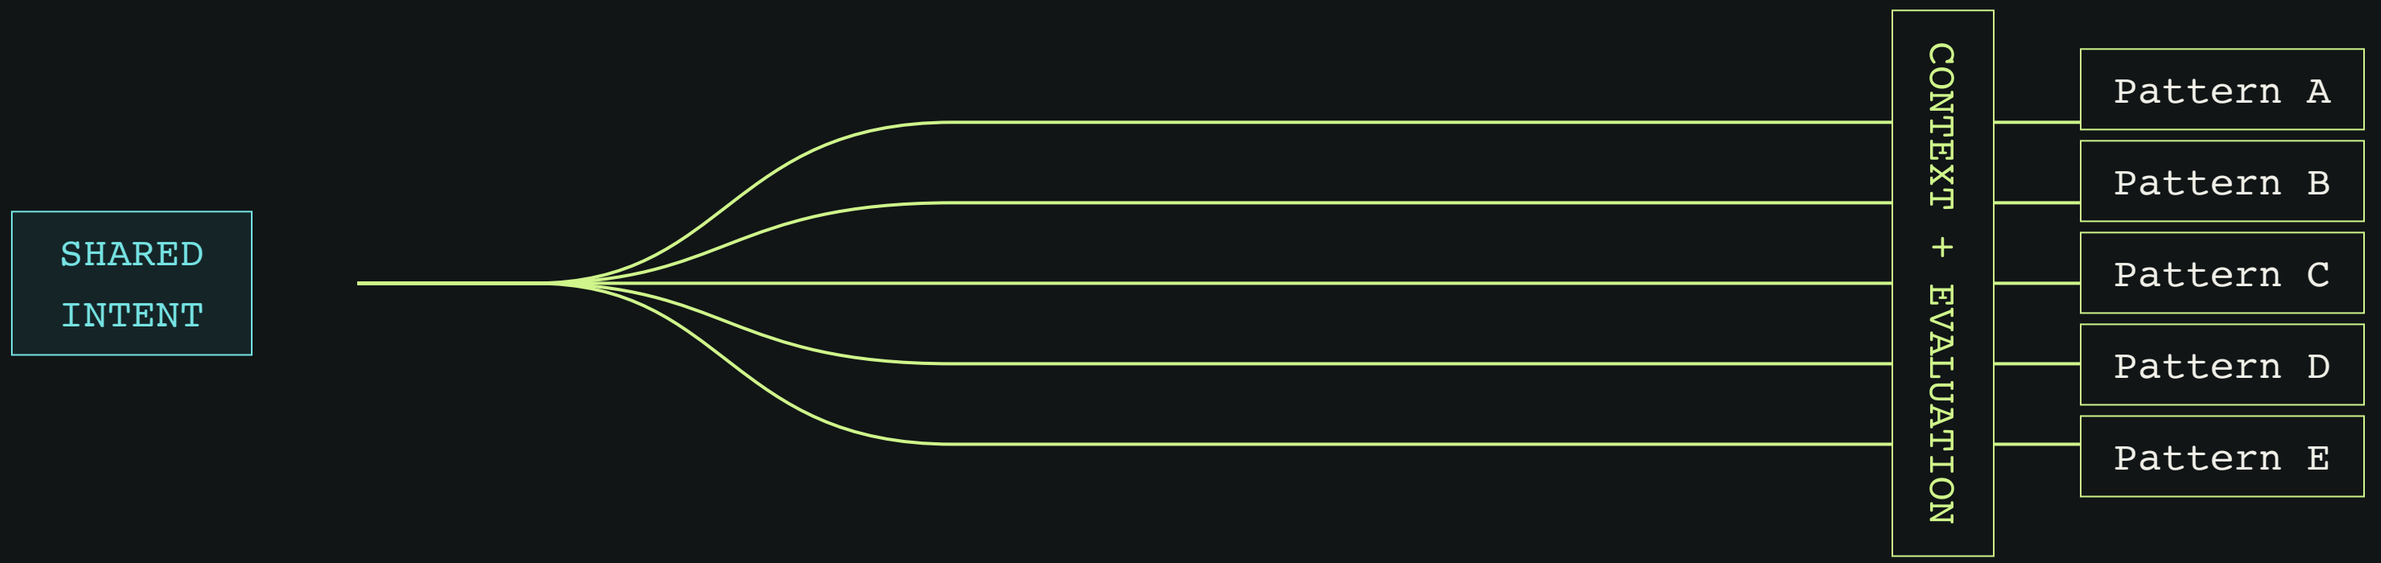
Julian Alvarado is a Senior Program Manager and Design Operations leader at Microsoft, helping Xbox Design scale through better systems, clearer processes, and responsible AI.

He works across teams and studios to improve how designers plan, learn, and deliver—building resilient, people-centered creative organizations.

Outside work, Julian is a daily CrossFit devotee and a “forever blue belt” in jiu-jitsu. A retired musician, he still loves live shows and enjoys drawing at home. Currently in rotation: Radiohead, Lenny Kravitz, Oasis, and Crosses.

AI doesn't just accelerate work. It accelerates variation.

More AI → More output → More variation → More decisions → More inconsistency



Distinct ideas. Shared constraints. Accountable decisions.

Three questions.
One coherent system.

GOVERN

Who can accept the choice?

Designer + engineer review

Product owner approves the rule

CONTEXT

What must the choice respect?

The user is starting alone

Invitations must be optional

EVALUATE

What makes it acceptable?

Complete setup without email

Check labels + errors + focus

Fix the form. Keep the lesson.

DESIGNER + ENGINEER + AI

- ✓ What the user needs
- ✓ What AI got wrong
- ✓ What the team changes
- ✓ Try the fix
- ✓ Save the lesson
- 06 Check the next result

The next proposed form lets the person start alone. A reviewer checks it again; saved guidance does not replace judgment.

THE NEXT RESULT / 06 — ILLUSTRATIVE

Can the next person start alone?

The next proposed form

Workspace name first. Invitations optional.

The review

Setup works without a teammate email.

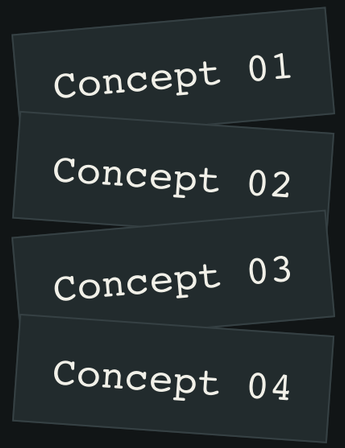
Still to check

Accessibility and whether it works for real users.

This requirement is met. The review is not over.

The next review: verify the behavior, not just the appearance.

Many possible outputs.
One accountable judgment.



Possibility is abundant.

SHARED QUALITY RUBRIC

- ☐ System compliance
- ☐ User needs
- ☐ Accessibility
- ☐ Brand
- ☐ Product strategy
- ☐ Engineering feasibility

Evidence + checks + expert critique

Human decision

Accept, revise,
or reject—with reasons.

OWNERSHIP STAYS HUMAN

Automated checks inform judgment. **They do not prove a design is good.**

Don't just save the output. Capture the learning.

Keep the decision and its reason.
Share them with the next team.

1. AI proposes a design — using shared guidance
2. The team reviews it — against the user's need
3. Save the lesson — decision + reason + owner
4. Update the guidance — for the next team + AI
5. Check the next result — do not assume success

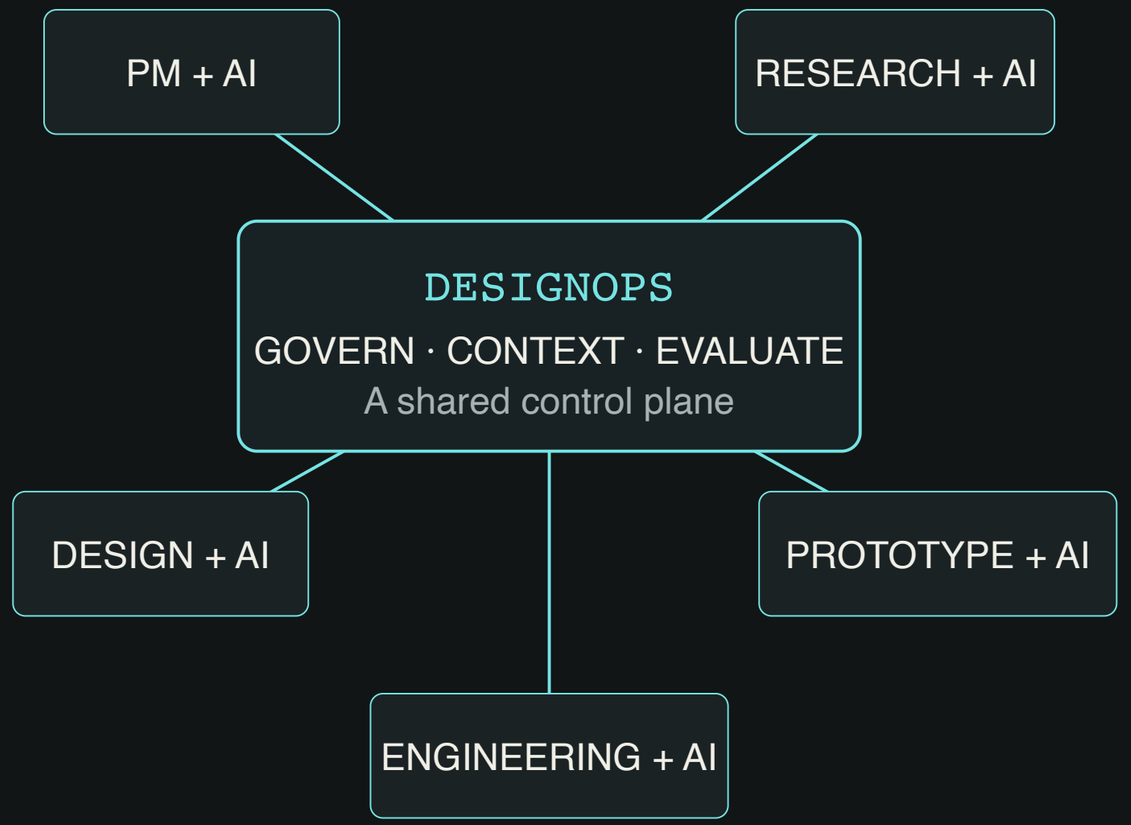
Same foundations.
A larger operating system.

PEOPLE
PROCESS
SYSTEMS

TRADITIONAL FOCUS	EXPANDED AI-ERA SCOPE
Optimize workflows	Architect human–AI workflows
Document processes	Encode organizational knowledge
Scale design systems	Make systems AI-readable
Enable designers	Enable product teams
Measure efficiency	Measure quality + outcomes
Facilitate consistency	Govern generated variation

Not doing all the work. Connecting the system.

DesignOps defines the conditions under which the work happens.



Domain experts still own product decisions.

Less waiting.
More learning, together.

Sequential handoffs

SIMPLIFIED BEFORE

Research → Brief → Mocks → Build → Release

Customer need → handoff → wait → handoff → wait → customer feedback

Shared intent. Joint execution.

PROPOSED PILOT

Mission · Milestones · Metrics → Spec ↔ concept ↔ HTML/CSS → Experiment ↔ learn

Customer need → PM + research + design + engineering → customer feedback

Measure production time, review effort, and learning cadence.

Hypothesis: production can shrink; experimentation can expand. Validate it.

Your next infrastructure is already in your documentation.

Guidelines	→	Machine-readable rules
Design system	→	AI-readable design system
Documentation	→	Reusable context packages
Critique	→	Evaluation rubrics

SHARED GUIDANCE / ILLUSTRATIVE

Lesson · solo setup

Invitations after setup

Reason · from the brief

The person may be starting alone.

Owner · product lead

Version 1, reviewed by the team

Review · after pilot / change

User test pending

DesignOps maintains the guidance.
The team checks it still applies.

Faster is not the whole score.

EARLY EXPERIMENTATION · APPROXIMATE VALUES

Brief → validated concept

Baseline	Pilot	Change
9d	6d	−3 d

median elapsed days

Design production rate

Baseline	Pilot	Change
~3—4/wk	~1—2/wk	≈−2/wk

accepted artifacts / in a team with 8w E2E Cycle

Change: midpoint estimate

Review effort

Baseline	Pilot	Change
4h	6h	+2 h

person-hours / accepted concept

Design-system compliance

Baseline	Pilot	Change
78%	90%	+12 pp

% applicable checks passed

Review effort rises. Is the trade-off worth it?

Compare equivalent work and verify user outcomes.

One workflow.
Thirty days. A measurable pilot.

DAYS 01—07

Map + identify

Find current AI use.
Locate inconsistency,
quality gaps, and risk.

DAYS 08—14

Select + package

Choose one workflow.
Package the relevant
organizational knowledge.

DAYS 15—21

Define + baseline

Name decision owners.
Set human checkpoints,
criteria, and baseline.

DAYS 22—30

Pilot + compare

Run equivalent work.
Review speed, quality,
and customer evidence.

Start with one team: design + engineering own the workflow; DesignOps facilitates.

If AI becomes part of how we design products,

who designs the system in which AI operates?

DesignOps.

AI doesn't eliminate the need for DesignOps.
It makes DesignOps infrastructure.

AI creates **capacity.**

Systems create **consistency.**

Stewardship creates **trust.**

DESIGNOPS AS AI INFRASTRUCTURE

Let's talk.

Where is your team generating faster than it can maintain coherence?

LinkedIn · alvaradolj1 ↗

GitHub · JulianAlvarado ↗

01 Who owns governance? ↗

02 Do we become an AI team? ↗

03 What is an AI-ready design system? ↗

04 What happens when models change? ↗

05 How do we prove value? ↗

06 How does the context stack fit? ↗

07 Where does this fit in the lifecycle? ↗

↖ BACK TO Q&A

Shared governance.
Explicit decision rights.

BOUNDARIES

Security + Legal

Data, privacy, risk, and policy.
Set non-negotiable constraints.

ACCESS

IT + platform teams

Approved technology and access.
Maintain secure capabilities.

OPERATING MODEL

DesignOps

Translate policy into workflows.
Facilitate context and evaluation.

PRODUCT JUDGMENT

Domain experts

PM, research, design, engineering.
Own quality and release decisions.

Good governance pre-approves safe paths—not every individual click.

↖ BACK TO Q&A

Architect the environment. Don't become the AI factory.

Governance
Rules + risk

Design intelligence
Consumable knowledge

Workflow architecture
With domain experts

Quality systems
Rubrics + evidence

Enablement
Playbooks + training

Measurement
Quality + outcomes

ENABLE DIFFERENT KINDS OF WORK

Prototype → Build → Improve → Grow → Maintain

Exploration, production, simplification, experimentation, and reliability need different controls.

↖ BACK TO Q&A

A component library with the reasoning attached.

KEEP THE FOUNDATION

- Tokens
- Components
- Patterns
- Guidelines

+

MAKE IT CONSUMABLE

- Semantic metadata
- Usage rules + constraints
- Examples + code equivalents
- Evaluation criteria

SHARED CONTEXT	Owner	Version	Source	Freshness	Access

Start with one pattern: its intent, code, valid uses, forbidden uses, and checks.

↖ BACK TO Q&A

Models change.
Your operating principles endure.

KNOWLEDGE	RULES	WORKFLOWS	EVALUATION

Model A ↔ Model B ↔ Model C

Re-run evaluation before switching.
Model-agnostic does not mean behavior-identical.

Canonical prompts encode domain knowledge—not one person's clever phrasing.

Co-author with PMs, researchers, designers, and engineers.
Version context, permissions, examples, and acceptance criteria together.

Measure the workflow.
Not the number of prompts.

Operational definitions for a scoped baseline and pilot

Metric	Definition / unit	Collection method
Brief → validated concept	Median elapsed days; fixed validation criteria	Brief timestamp → review + user evidence
Design production rate	Accepted artifacts per team-week	Same artifact class, scope, and acceptance rubric
Production / review effort	Person-hours per accepted concept; report separately	Time log including critique + validation
Rework / revision rate	% reviewed artifacts requiring a major revision	Revision log; agree “major” before starting
System compliance	% applicable design-system checks passed	Token, component, and interaction-pattern audit
Accessibility / quality defects	Open defects per 10 reviewed screens + severity	Manual and automated checks, deduplicated
Design-to-code divergence	% audited UI instances violating approved intent	Paired design/code review of a defined sample
Experimentation / reuse	Completed experiments/week; % work citing approved context	Decision log; context provenance in artifacts

Match scope + staffing. Record sample size + spread. Audit sampling bias. Track customer task success separately.
Small pilots are directional—not causal proof. Set adoption / revise / stop criteria before seeing results.

↖ BACK TO Q&A

The advantage isn't the model. It's the memory.

What does your organization allow the model to know?

05 **DesignOps**
Connects the system · enables adoption

04 **Stewardship**
Guardrails · exclusions · human checkpoints

03 **Organizational context**
Tokens · components · research · brand · accessibility
Terminology · workflows · repos · product experiences

02 **Agents + tools**
Design assistants · coding agents · research tools

01 **Models**
Replaceable capabilities, not institutional knowledge

↖ BACK TO Q&A

Same infrastructure.
Every stage of the journey.

01	02	03	04	05
DISCOVER	DEFINE	DESIGN	BUILD	MEASURE
Synthesize research	Frame the problem	Explore concepts	Generate components	Interpret outcomes

SHARED INFRASTRUCTURE — GOVERN — CONTEXT — EVALUATE